

# drift

October 16, 2018

## 0.1 Dealing with Drift

We will be attempting to see if intelligence and higher income have any correlation. Let us start with our imports and loading our dataset.

```
In [1]: import pandas as pd
```

```
In [2]: incomeintel = pd.read_csv('IncomeIntel.txt')
```

```
In [3]: incomeintel
```

```
Out[3]:
```

	grad_year	gre_qnt	salary_p4
0	2001.0	739.737072	67400.475185
1	2001.0	721.811673	67600.584142
2	2001.0	736.277908	58704.880589
3	2001.0	770.498485	64707.290345
4	2001.0	735.002861	51737.324165
5	2001.0	763.876037	64010.822579
6	2001.0	738.758659	60080.107481
7	2001.0	706.407471	56263.309815
8	2001.0	705.886037	62109.859243
9	2001.0	700.971986	50189.704747
10	2001.0	709.754522	58721.753127
11	2001.0	734.854582	65380.594586
12	2001.0	753.384151	52857.212365
13	2001.0	690.312090	63572.217765
14	2001.0	774.154371	65892.177035
15	2001.0	726.377225	67454.545201
16	2001.0	702.735945	59346.670232
17	2001.0	723.806542	70031.012603
18	2001.0	758.051159	53441.672888
19	2001.0	711.063082	61008.652046
20	2001.0	702.975969	50065.932451
21	2001.0	733.877837	75612.225369
22	2001.0	735.918767	59580.620375
23	2001.0	749.069115	57825.611782
24	2001.0	732.581793	52809.225854
25	2001.0	728.050446	57492.084316
26	2001.0	690.265988	64686.224351

27	2001.0	732.448836	53067.021394
28	2001.0	724.755887	58902.707320
29	2001.0	721.739038	62094.061567
..	...	...	...
970	2013.0	158.578197	79263.470892
971	2013.0	147.667305	104782.627567
972	2013.0	160.086274	94013.946074
973	2013.0	156.289493	74032.543183
974	2013.0	150.340044	84220.290724
975	2013.0	163.054596	74940.546965
976	2013.0	157.624151	83293.343135
977	2013.0	150.927266	78340.908128
978	2013.0	157.393763	91066.889575
979	2013.0	154.449630	87169.012509
980	2013.0	153.756644	90033.601423
981	2013.0	150.796371	98650.768576
982	2013.0	150.691700	70455.885421
983	2013.0	153.639896	91133.301177
984	2013.0	150.374470	91796.617819
985	2013.0	162.350725	73780.832249
986	2013.0	155.803279	96927.925237
987	2013.0	159.111662	71875.246552
988	2013.0	158.338350	103357.966587
989	2013.0	162.308518	73780.472319
990	2013.0	156.651125	79055.571295
991	2013.0	153.836045	91529.313046
992	2013.0	149.542467	75940.200168
993	2013.0	155.349020	97688.397380
994	2013.0	161.767399	75260.194609
995	2013.0	160.441025	100430.166532
996	2013.0	160.431891	82198.200872
997	2013.0	154.254526	84340.214218
998	2013.0	162.036321	87600.881985
999	2013.0	156.946735	82854.576903

[1000 rows x 3 columns]

Now that we have our data loaded up let us use the Statsmodel for our OLS model. It allows for more information about the model unlike our previously used scikit-learn model.

```
In [4]: import statsmodels.api as sm

y, X = incomeintel['salary_p4'], incomeintel['gre_qnt']
X = sm.add_constant(X, prepend=False)

ols = sm.OLS(y, X)
ols_result = ols.fit()
ols_result.summary()
```

```

Out[4]: <class 'statsmodels.iolib.summary.Summary'>
"""
                                OLS Regression Results
=====
Dep. Variable:                  salary_p4    R-squared:                        0.263
Model:                            OLS        Adj. R-squared:                   0.262
Method:                    Least Squares    F-statistic:                       356.3
Date:                Tue, 16 Oct 2018    Prob (F-statistic):                3.43e-68
Time:                20:45:05    Log-Likelihood:                    -10673.
No. Observations:                1000    AIC:                             2.135e+04
Df Residuals:                    998    BIC:                             2.136e+04
Df Model:                        1
Covariance Type:                nonrobust
=====
                                coef    std err          t      P>|t|      [0.025      0.975]
-----
gre_qnt         -25.7632         1.365    -18.875      0.000     -28.442     -23.085
const          8.954e+04       878.764     101.895      0.000      8.78e+04      9.13e+04
=====
Omnibus:                 9.118    Durbin-Watson:                   1.424
Prob(Omnibus):            0.010    Jarque-Bera (JB):                 9.100
Skew:                    0.230    Prob(JB):                         0.0106
Kurtosis:                3.077    Cond. No.                      1.71e+03
=====

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 1.71e+03. This might indicate that there are
strong multicollinearity or other numerical problems.
"""

```

The coefficient is 103.0949 and the standard error is 1.733. We can see the rest of the results of the OLS when we run the `ols.summary()` function. The result is significant. The coefficient is -25.7632 and the standard error is 1.365. We will now plot the data on a scatterplot. We'll be using the standard scatter method which pandas provides us.

```


In [5]: incomeintel.plot.scatter(x='grad_year', y='gre_qnt')

```

```

Out[5]: <matplotlib.axes._subplots.AxesSubplot at 0x1144a1cf8>

```




We can see a clear problem come up - in 2011 the scoring system for the GRE changed. A score of 750 doesn't even exist post 2011! For us to have any kind of meaningful analysis we must scale these values.

The ideal way to scale these values would be to have the percentile values for each of the observations, and report all the scores as percentiles. The GRE website also gives us a conversion rate which we can use, which should also give us acceptable results. What we will resort to is simply scale the values, by multiplying all the values post 2011 with 800/170. We would like to stress that this is not an ideal approach, but will suffice to go through this experiment.

```
In [6]: incomeintel.loc[incomeintel.grad_year > 2011, 'gre_qnt'] = incomeintel.g
```

```
In [7]: incomeintel.plot.scatter(x='grad_year', y='salary_p4')
```

```
Out[7]: <matplotlib.axes._subplots.AxesSubplot at 0x1172dfe48>
```



We can see a trend here - with increasing year, there is an increase in the salary. This might effect our results, and we would need to detrend our data. Because it is not panel data, we cannot use differencing or log differencing methods. We can attempt to fix this instead by finding the average increase per year to find the average growth rate. We then adjust by dividing all the salaries by  $(1 + \text{avg\_growth\_rate}) ** (\text{grad\_year} - 2001)$ . This will not affect the first year, but adjust the rest of the years. We add a new column called `adjusted_salary` to store these new values.


```
In [8]: avg_inc_by_year = incomeintel['salary_p4'].groupby(incomeintel['grad_year']).mean().values
```

```
In [9]: avg_growth_rate = ((avg_inc_by_year[1:] - avg_inc_by_year[:-1]) / avg_inc_by_year[:-1]).values
```

```
In [10]: incomeintel['adjusted_salary'] = incomeintel['salary_p4'] / ((1 + avg_growth_rate) ** (grad_year - 2001))
```

```
In [11]: incomeintel.plot.scatter(x='grad_year', y='adjusted_salary')
```

```
Out[11]: <matplotlib.axes._subplots.AxesSubplot at 0x11733dcc0>
```



We've now adjusted our salary, and also adjusted for the GRE scores which have changed. Let us now try and train another model which can better represent our question!

```
In [12]: y, X = incomeintel['adjusted_salary'], incomeintel['gre_qnt']
```

```
X = sm.add_constant(X, prepend=False)
```

```
newols = sm.OLS(y, X)
newols_result = newols.fit()
newols_result.summary()
```

```
Out[12]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                                OLS Regression Results
=====
Dep. Variable:          adjusted_salary    R-squared:                0.000
Model:                  OLS              Adj. R-squared:           -0.001
Method:                 Least Squares     F-statistic:              0.1814
Date:                  Tue, 16 Oct 2018   Prob (F-statistic):       0.670
Time:                  20:45:06          Log-Likelihood:           -10291.
No. Observations:      1000             AIC:                     2.059e+04
Df Residuals:          998              BIC:                     2.060e+04
Df Model:              1
Covariance Type:       nonrobust
=====
                                coef    std err          t      P>|t|      [0.025    0.975]
=====
```

```

-----
gre_qnt      -0.6217      1.460      -0.426      0.670      -3.486      2.242
const       6.185e+04    1024.106     60.390      0.000      5.98e+04    6.39e+04
=====
Omnibus:                                0.758    Durbin-Watson:                        2.026
Prob(Omnibus):                          0.685    Jarque-Bera (JB):                      0.672
Skew:                                    0.059    Prob(JB):                              0.715
Kurtosis:                               3.046    Cond. No.                             3.18e+03
=====

```

Warnings:

```

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 3.18e+03. This might indicate that there are
strong multicollinearity or other numerical problems.
"""

```

After we adjust for time drift and fix the GRE scores, we get a coefficient of -0.6217, and standard error of 1.460, and we can see that the result is not significant. This means that GRE quant scores doesn't really account for your salary. This is a very important example of how we need to carefully examine our data and the nature of our data before we proceed with any analysis!